

PageRank Project Report

Zecheng Xu (zx78) Shuo Xiong (sx78) Qianying Zhang (qz226)

We implement node-by-node and blocked PageRank algorithm, as well as Jacobi version Blocked PageRank and Random partition Blocked Pagerank

Filter Parameter:

netid: zx78

rejectMin: $0.87 * 0.99 = 0.8613$

rejectLimit: $0.87 * 0.99 + 0.01 = 0.8713$

number of edges selected : 7524770

We preprocessed the data (DataProcessor.java) from edges.txt and create a new file that used as input for MapReduce. The format of each line of the new file is:

< NodeID + PageRank + Degree + NeighborNodes >

All input file, output files and .jar is stored in bucket:
edu-cornell-cs-cs5300s14-zx78-project2

Simple PageRank

In simplePageRank package, there are four java classes:

- 1) SimplePageRank
- 2) SimplePageRankMapper
- 3) SimplePageRankReducer
- 4) Counter

SimplePageRank

This class contains the main method. MapReduce job is configured with certain input/output format and input/output path.

In our implementation, we stored intermediate results in HDFS to improve efficiency since it's faster than read from/write to AWS S3.

SimplePageRankMapper

Since we use the default *TextInputFormat*, the key of the InputSplit is the offset of line, value is the content of the line. In our new created input file, each line correspond to a node.

We pass two kinds of information to reducer:

1. <srcNode, "prevNodeInfo" + PageRank + srcDegree> if srcDegree = 0

<srcNode, "prevNodeInfo" + PageRank + srcDegree + dstNodes>, otherwise this is for the purpose of reconstructing the node.

2. <dstNode, srcPageRank/srcDegree>

this is to compute new pagerank value for "dstNode".

SimplePageRankReducer

In reducer, it simply reconstruct the node and using the second type of information passed from mapper to compute new PageRank value for this node.

Then it increments the RESIDUAL_COUNTER in Counter class. The output of reducer in current iteration is the input for mapper in next iteration.

Counter

Since we use counter, in each pass, only one MapReduce job is required. There is only one field in Counter: RESIDUAL_COUNTER.

Every reducer increments RESIDUAL_COUNTER after computing residual error.

After all reducer finishes working, we output the average residual error according to the counter value.

Result of Simple PageRank

run 6 iterations:

```
Iteration 0 Avg. residual error: 2.3385827422062664
Iteration 1 Avg. residual error: 0.3229186718255184
Iteration 2 Avg. residual error: 0.19205372241145308
Iteration 3 Avg. residual error: 0.09416565085445763
Iteration 4 Avg. residual error: 0.06287620062460779
Iteration 5 Avg. residual error: 0.03389726990207667
```

Blocked PageRank

It takes 7 passes to converge globally

There are four classes in this package, code structure is similar to simple PageRank:

- 1) BlockedPageRank
- 2) BlockedPageRankMapper
- 3) BlockedPageRankReducer
- 4) BlockCounter

BlockedPageRank

The class contains main method. Block boundaries are hard-coded for convenience.

MapReduce passes are implemented until global average residual error is lower than desired lower bound: 0.001. In the end of each termination, we output the average iteration within blocks and the average residual error. After termination, output pageranks of highest numbered nodes in every blocks.

BlockedPageRankMapper

In mapper, we implement the *blockIDofNode* method and runs in constant time.

Three types of information were passed to reducer:

1. If $\text{srcDegree} = 0$:
 $\langle \text{srcBlock}, \text{"prevNodeInfo"} + \text{srcNode} + \text{pageRank} + \text{srcDegree} \rangle$
otherwise:
 $\langle \text{srcBlock}, \text{"prevNodeInfo"} + \text{srcNode} + \text{PageRank} + \text{srcDegree} + \text{dstNodes} \rangle$
2. for edges within block, $\{ \langle u, v \rangle \mid u \in B \wedge u \rightarrow v \}$
 $\langle \text{dstBlock}, \text{"BE"} + \text{srcNode}(u) + \text{dstNode}(v) \rangle$
3. for edges entering $\{ \langle u, v, R \rangle \mid u \notin B \wedge v \in B \wedge u \rightarrow v \wedge R = \text{PR}(u)/\text{deg}(u) \}$
 $\langle \text{dstBlock}, \text{"BC"} + \text{dstNode}(v) + \text{srcPR}(u)/\text{deg}(u) \rangle$

BlockedPageRankReducer

We use several HashMaps to store node information and edges information.

In *IteratorBlockOnce*, we first backup current pageranks and then recompute pageranks according to edges related and then compute average residual errors within the block.

The termination condition in reducer is while the nodes within the block are converged.

After termination, increase counter REDUCER_COUNTER by 1,

INNER_ITERATION_COUNTER by # of iterations of *IteratorBlockOnce*,

RESIDUAL_COUNTER by the total residual errors in the block and corresponding

PAGERANK[blockID] by the pagerank of the highest numbered nodes in block.

BlockCounter:

RESIDUAL_COUNTER: sum of residual errors

REDUCER_COUNTER: # of reducers

INNER_ITERATION_COUNTER: sum of iterations within blocks

$$\text{Avg iterations} = \text{INNER_ITERATION_COUNTER} / \text{REDUCER_COUNTER}$$

PAGERANK0 – PAGERANK67: store pagerank of highest numbered nodes in the corresponding block

Result of Blocked PageRank:

[pass 0] Avg. residual error: 2.8157218713967214, Avg. iteration: 17.485294
[pass 1] Avg. residual error: 0.03790001912624106, Avg. iteration: 7.161765
[pass 2] Avg. residual error: 0.023992188521758957, Avg. iteration: 5.867647
[pass 3] Avg. residual error: 0.009920314455334589, Avg. iteration: 3.897059
[pass 4] Avg. residual error: 0.00395191477851447, Avg. iteration: 2.558824
[pass 5] Avg. residual error: 0.0010555445699314391, Avg. iteration: 1.426471
[pass 6] Avg. residual error: 6.622697747534988E-4, Avg. iteration: 1.176471

sample pageranks:

[block 0] Max node: 10327 PageRank: 1.87285E-6
[block 1] Max node: 20372 PageRank: 5.2155E-7
[block 2] Max node: 30628 PageRank: 3.0675E-7
[block 3] Max node: 40644 PageRank: 2.8092E-7
[block 4] Max node: 50461 PageRank: 3.0002E-7
[block 5] Max node: 60840 PageRank: 2.189E-7
[block 6] Max node: 70590 PageRank: 3.0011E-7
[block 7] Max node: 80117 PageRank: 2.189E-7
[block 8] Max node: 90496 PageRank: 8.664012E-4
[block 9] Max node: 100500 PageRank: 2.5982E-7
[block 10] Max node: 110566 PageRank: 2.02975E-6
[block 11] Max node: 120944 PageRank: 4.8272E-7
[block 12] Max node: 130998 PageRank: 2.3553E-7
[block 13] Max node: 140573 PageRank: 6.2776E-7
[block 14] Max node: 150952 PageRank: 2.189E-7
[block 15] Max node: 161331 PageRank: 2.189E-7
[block 16] Max node: 171153 PageRank: 5.6226E-7
[block 17] Max node: 181513 PageRank: 3.4393E-7
[block 18] Max node: 191624 PageRank: 5.27754E-6
[block 19] Max node: 202003 PageRank: 4.14836E-6
[block 20] Max node: 212382 PageRank: 2.6175E-7
[block 21] Max node: 222761 PageRank: 0.0012286764
[block 22] Max node: 232592 PageRank: 1.2014965E-4
[block 23] Max node: 242877 PageRank: 5.6451E-7
[block 24] Max node: 252937 PageRank: 2.189E-7
[block 25] Max node: 263148 PageRank: 2.9941E-7
[block 26] Max node: 273209 PageRank: 2.189E-7
[block 27] Max node: 283472 PageRank: 3.658E-7
[block 28] Max node: 293254 PageRank: 2.6175E-7
[block 29] Max node: 303042 PageRank: 4.94703E-6

[block 30]	Max node: 313369	PageRank: 8.7558E-7
[block 31]	Max node: 323521	PageRank: 2.189E-7
[block 32]	Max node: 333882	PageRank: 5.1807E-7
[block 33]	Max node: 343662	PageRank: 2.5204E-7
[block 34]	Max node: 353644	PageRank: 4.3791E-7
[block 35]	Max node: 363928	PageRank: 3.5619E-7
[block 36]	Max node: 374235	PageRank: 3.26453E-6
[block 37]	Max node: 384553	PageRank: 4.272E-7
[block 38]	Max node: 394928	PageRank: 6.6669E-7
[block 39]	Max node: 404711	PageRank: 2.4565E-7
[block 40]	Max node: 414616	PageRank: 3.1166E-7
[block 41]	Max node: 424746	PageRank: 2.02404E-6
[block 42]	Max node: 434706	PageRank: 2.81297E-6
[block 43]	Max node: 444488	PageRank: 3.82919E-6
[block 44]	Max node: 454284	PageRank: 4.9751E-7
[block 45]	Max node: 464397	PageRank: 1.105941E-5
[block 46]	Max node: 474195	PageRank: 9.6575E-7
[block 47]	Max node: 484049	PageRank: 5.9304E-7
[block 48]	Max node: 493967	PageRank: 2.189E-7
[block 49]	Max node: 503751	PageRank: 6.68699E-6
[block 50]	Max node: 514130	PageRank: 6.1659E-7
[block 51]	Max node: 524509	PageRank: 9.8640168E-4
[block 52]	Max node: 534708	PageRank: 2.300763E-5
[block 53]	Max node: 545087	PageRank: 0.00257605008
[block 54]	Max node: 555466	PageRank: 0.00120386992
[block 55]	Max node: 565845	PageRank: 1.64321E-6
[block 56]	Max node: 576224	PageRank: 1.10812E-6
[block 57]	Max node: 586603	PageRank: 9.565E-7
[block 58]	Max node: 596584	PageRank: 4.10833E-6
[block 59]	Max node: 606366	PageRank: 4.2766E-7
[block 60]	Max node: 616147	PageRank: 6.2093E-7
[block 61]	Max node: 626447	PageRank: 1.565152E-5
[block 62]	Max node: 636239	PageRank: 1.02763E-6
[block 63]	Max node: 646021	PageRank: 3.1193E-7
[block 64]	Max node: 655803	PageRank: 2.189E-7
[block 65]	Max node: 665665	PageRank: 9.9855E-7
[block 66]	Max node: 675447	PageRank: 1.0857E-6
[block 67]	Max node: 685229	PageRank: 3.5609E-7

Randomed Block Partition

We define $\text{hash}(\text{nodeID}) = \text{nodeID}$

so that $\text{blockIDofNode}(\text{nodeID}) = \text{nodeID} \% \text{NUM_OF_BLOCKS}$

It takes 22 passes to converge globally, much slower than intelligently partitioned blocks since in bad partition, most edges are between blocks and few edges are within blocks. The performance of random block partition is basically differs little from node-by-node implementation.

Result:

[pass 0] Avg. residual error: 2.3392050529342736, Avg. iteration: 3.000000
[pass 1] Avg. residual error: 0.322338531099608, Avg. iteration: 2.720588
[pass 2] Avg. residual error: 0.19120350060835012, Avg. iteration: 2.000000
[pass 3] Avg. residual error: 0.09352085277190784, Avg. iteration: 2.000000
[pass 4] Avg. residual error: 0.062049689719819, Avg. iteration: 2.000000
[pass 5] Avg. residual error: 0.03343268009401611, Avg. iteration: 2.000000
[pass 6] Avg. residual error: 0.026737168901070005, Avg. iteration: 2.000000
[pass 7] Avg. residual error: 0.016329769738696306, Avg. iteration: 2.000000
[pass 8] Avg. residual error: 0.014039672567498987, Avg. iteration: 2.000000
[pass 9] Avg. residual error: 0.009590845800121406, Avg. iteration: 2.000000
[pass 10] Avg. residual error: 0.008232902427247786, Avg. iteration: 2.000000
[pass 11] Avg. residual error: 0.005990730327351765, Avg. iteration: 2.000000
[pass 12] Avg. residual error: 0.0052024045009531105, Avg. iteration: 2.000000
[pass 13] Avg. residual error: 0.003905294158274054, Avg. iteration: 2.000000
[pass 14] Avg. residual error: 0.003378213739316799, Avg. iteration: 2.000000
[pass 15] Avg. residual error: 0.002622637238369467, Avg. iteration: 2.000000
[pass 16] Avg. residual error: 0.0022454500814906527, Avg. iteration: 2.000000
[pass 17] Avg. residual error: 0.0017787145639380354, Avg. iteration: 2.000000
[pass 18] Avg. residual error: 0.0015187079725565138, Avg. iteration: 2.000000
[pass 19] Avg. residual error: 0.0012179097482476395, Avg. iteration: 2.000000
[pass 20] Avg. residual error: 0.001034286744685084, Avg. iteration: 1.779412
[pass 21] Avg. residual error: 8.398437107079812E-4, Avg. iteration: 1.000000

Sample pageranks:

[block 0] Max node: 685168 PageRank: 3.49887E-6
[block 1] Max node: 685169 PageRank: 7.26665E-6
[block 2] Max node: 685170 PageRank: 7.36389E-6
[block 3] Max node: 685171 PageRank: 7.84502E-6
[block 4] Max node: 685172 PageRank: 3.5926E-7
[block 5] Max node: 685173 PageRank: 4.2094E-7

[block 6]	Max node: 685174	PageRank: 6.71932E-6
[block 7]	Max node: 685175	PageRank: 8.16817E-6
[block 8]	Max node: 685176	PageRank: 2.8461E-7
[block 9]	Max node: 685177	PageRank: 6.71489E-6
[block 10]	Max node: 685178	PageRank: 6.68529E-6
[block 11]	Max node: 685179	PageRank: 6.71046E-6
[block 12]	Max node: 685180	PageRank: 5.9564E-7
[block 13]	Max node: 685181	PageRank: 3.5831E-7
[block 14]	Max node: 685182	PageRank: 5.3393E-7
[block 15]	Max node: 685183	PageRank: 4.7868E-7
[block 16]	Max node: 685184	PageRank: 4.8401E-7
[block 17]	Max node: 685185	PageRank: 3.5831E-7
[block 18]	Max node: 685186	PageRank: 6.8677E-7
[block 19]	Max node: 685187	PageRank: 3.7121E-7
[block 20]	Max node: 685188	PageRank: 4.8401E-7
[block 21]	Max node: 685189	PageRank: 4.7868E-7
[block 22]	Max node: 685190	PageRank: 4.224E-7
[block 23]	Max node: 685191	PageRank: 7.815E-7
[block 24]	Max node: 685192	PageRank: 4.8401E-7
[block 25]	Max node: 685193	PageRank: 4.8401E-7
[block 26]	Max node: 685194	PageRank: 3.5831E-7
[block 27]	Max node: 685195	PageRank: 4.7868E-7
[block 28]	Max node: 685196	PageRank: 4.5694E-7
[block 29]	Max node: 685197	PageRank: 3.2065E-7
[block 30]	Max node: 685198	PageRank: 4.8401E-7
[block 31]	Max node: 685199	PageRank: 1.45206E-6
[block 32]	Max node: 685200	PageRank: 4.6583E-7
[block 33]	Max node: 685201	PageRank: 4.6583E-7
[block 34]	Max node: 685202	PageRank: 6.5998E-7
[block 35]	Max node: 685203	PageRank: 4.2232E-7
[block 36]	Max node: 685204	PageRank: 4.8401E-7
[block 37]	Max node: 685205	PageRank: 3.5831E-7
[block 38]	Max node: 685206	PageRank: 3.7121E-7
[block 39]	Max node: 685207	PageRank: 6.8677E-7
[block 40]	Max node: 685208	PageRank: 4.8401E-7
[block 41]	Max node: 685209	PageRank: 4.8401E-7
[block 42]	Max node: 685210	PageRank: 4.8401E-7
[block 43]	Max node: 685211	PageRank: 3.5831E-7
[block 44]	Max node: 685212	PageRank: 1.79671E-6
[block 45]	Max node: 685213	PageRank: 1.79671E-6
[block 46]	Max node: 685214	PageRank: 1.79671E-6

[block 47]	Max node: 685215	PageRank: 1.79671E-6
[block 48]	Max node: 685216	PageRank: 1.79671E-6
[block 49]	Max node: 685217	PageRank: 1.79671E-6
[block 50]	Max node: 685218	PageRank: 4.7868E-7
[block 51]	Max node: 685219	PageRank: 2.189E-7
[block 52]	Max node: 685220	PageRank: 8.1196E-7
[block 53]	Max node: 685221	PageRank: 6.5998E-7
[block 54]	Max node: 685222	PageRank: 7.8008E-7
[block 55]	Max node: 685223	PageRank: 3.5831E-7
[block 56]	Max node: 685224	PageRank: 3.7482E-7
[block 57]	Max node: 685225	PageRank: 4.2232E-7
[block 58]	Max node: 685226	PageRank: 4.2232E-7
[block 59]	Max node: 685227	PageRank: 1.9606E-6
[block 60]	Max node: 685228	PageRank: 1.88119E-6
[block 61]	Max node: 685229	PageRank: 3.5831E-7
[block 62]	Max node: 685162	PageRank: 2.72815E-6
[block 63]	Max node: 685163	PageRank: 2.72815E-6
[block 64]	Max node: 685164	PageRank: 2.72372E-6
[block 65]	Max node: 685165	PageRank: 2.69502E-6
[block 66]	Max node: 685166	PageRank: 2.71929E-6
[block 67]	Max node: 685167	PageRank: 2.72815E-6

Gauss-Seidel

The implementation is similar to that of Blocked PageRank, only with a little difference in the reducer.

The following equation defines the PageRank value using Gauss-Seidel method:

$$PR(i)^{(k+1)} = (1 - \alpha) / NUM_OF_NODES + \alpha \left(\sum_{i < j} PR(j)^k / deg(j) + \sum_{i > j} PR(j)^{(k+1)} / deg(j) \right)$$

$< j, i > \in E$

The number of passes using Gauss-Seidel method is the same as Jacobi method, however, the average number of iterations in reducer is smaller than Jacobi.

Result:

[pass 0]	Avg. residual error: 2.816151834165567,	Avg. iteration: 10.088235
[pass 1]	Avg. residual error: 0.038937140946554455,	Avg. iteration: 5.102941
[pass 2]	Avg. residual error: 0.02530163499966147,	Avg. iteration: 4.367647

[pass 3] Avg. residual error: 0.01116037507638657, Avg. iteration: 3.220588
[pass 4] Avg. residual error: 0.005083253142829371, Avg. iteration: 2.367647
[pass 5] Avg. residual error: 0.0019187138620793164, Avg. iteration: 1.661765
[pass 6] Avg. residual error: 8.445789655928958E-4, Avg. iteration: 1.308824

Sample PageRanks:

[block 0] Max node: 10327 PageRank: 1.87275E-6
[block 1] Max node: 20372 PageRank: 5.2159E-7
[block 2] Max node: 30628 PageRank: 3.062E-7
[block 3] Max node: 40644 PageRank: 2.8092E-7
[block 4] Max node: 50461 PageRank: 2.9947E-7
[block 5] Max node: 60840 PageRank: 2.189E-7
[block 6] Max node: 70590 PageRank: 3.0011E-7
[block 7] Max node: 80117 PageRank: 2.189E-7
[block 8] Max node: 90496 PageRank: 8.6470304E-4
[block 9] Max node: 100500 PageRank: 2.5989E-7
[block 10] Max node: 110566 PageRank: 2.03005E-6
[block 11] Max node: 120944 PageRank: 4.7697E-7
[block 12] Max node: 130998 PageRank: 2.3454E-7
[block 13] Max node: 140573 PageRank: 6.3159E-7
[block 14] Max node: 150952 PageRank: 2.189E-7
[block 15] Max node: 161331 PageRank: 2.189E-7
[block 16] Max node: 171153 PageRank: 5.6224E-7
[block 17] Max node: 181513 PageRank: 3.4387E-7
[block 18] Max node: 191624 PageRank: 5.27752E-6
[block 19] Max node: 202003 PageRank: 4.14219E-6
[block 20] Max node: 212382 PageRank: 2.6175E-7
[block 21] Max node: 222761 PageRank: 0.00120921152
[block 22] Max node: 232592 PageRank: 1.0427226E-4
[block 23] Max node: 242877 PageRank: 5.6716E-7
[block 24] Max node: 252937 PageRank: 2.189E-7
[block 25] Max node: 263148 PageRank: 2.9941E-7
[block 26] Max node: 273209 PageRank: 2.189E-7
[block 27] Max node: 283472 PageRank: 3.657E-7
[block 28] Max node: 293254 PageRank: 2.6175E-7
[block 29] Max node: 303042 PageRank: 4.94705E-6
[block 30] Max node: 313369 PageRank: 8.9347E-7
[block 31] Max node: 323521 PageRank: 2.189E-7
[block 32] Max node: 333882 PageRank: 5.1807E-7
[block 33] Max node: 343662 PageRank: 2.5204E-7
[block 34] Max node: 353644 PageRank: 4.5195E-7

[block 35]	Max node: 363928	PageRank: 3.564E-7
[block 36]	Max node: 374235	PageRank: 3.76355E-6
[block 37]	Max node: 384553	PageRank: 4.2771E-7
[block 38]	Max node: 394928	PageRank: 6.6675E-7
[block 39]	Max node: 404711	PageRank: 2.4567E-7
[block 40]	Max node: 414616	PageRank: 3.1167E-7
[block 41]	Max node: 424746	PageRank: 2.1234E-6
[block 42]	Max node: 434706	PageRank: 2.81472E-6
[block 43]	Max node: 444488	PageRank: 3.80724E-6
[block 44]	Max node: 454284	PageRank: 4.9751E-7
[block 45]	Max node: 464397	PageRank: 1.117764E-5
[block 46]	Max node: 474195	PageRank: 9.659E-7
[block 47]	Max node: 484049	PageRank: 5.9322E-7
[block 48]	Max node: 493967	PageRank: 2.189E-7
[block 49]	Max node: 503751	PageRank: 6.7826E-6
[block 50]	Max node: 514130	PageRank: 6.1659E-7
[block 51]	Max node: 524509	PageRank: 9.8739784E-4
[block 52]	Max node: 534708	PageRank: 2.281982E-5
[block 53]	Max node: 545087	PageRank: 0.00258417872
[block 54]	Max node: 555466	PageRank: 0.00120476768
[block 55]	Max node: 565845	PageRank: 1.64295E-6
[block 56]	Max node: 576224	PageRank: 1.10794E-6
[block 57]	Max node: 586603	PageRank: 9.5652E-7
[block 58]	Max node: 596584	PageRank: 4.07348E-6
[block 59]	Max node: 606366	PageRank: 4.2702E-7
[block 60]	Max node: 616147	PageRank: 6.2176E-7
[block 61]	Max node: 626447	PageRank: 1.716388E-5
[block 62]	Max node: 636239	PageRank: 1.02692E-6
[block 63]	Max node: 646021	PageRank: 3.1193E-7
[block 64]	Max node: 655803	PageRank: 2.189E-7
[block 65]	Max node: 665665	PageRank: 9.9978E-7
[block 66]	Max node: 675447	PageRank: 1.08566E-6
[block 67]	Max node: 685229	PageRank: 3.568E-7